

Assignment Project Exam Help

Week 3: Mutual Exclusion Principles, Parallel Performance, Spin Locks and Contention

<https://powcoder.com>

MPCS 52060: Parallel Programming

University of Chicago

Add WeChat powcoder

Assignment Project Exam Help

Principles of Concurrent
Programs

<https://powcoder.com>

Add WeChat powcoder

Two types of formal properties in asynchronous computation:

- **Safety** Properties:

- Nothing bad happens ever

- For example - a traffic light never displays green in all directions, even if power fails.

- **Liveness** Properties:

- States that a particular “good” thing will happen.

- For example - a red light will eventually turn green.

Formalizing Critical Sections

Synchronization primitives (e.g., locks) need to adhere to the following properties and principles about critical sections in order to be correct:

- **Mutual Exclusion Property:**

- Critical sections of different threads do not overlap. Only one thread is executing a critical section at a time

Guarantees that a computation's results are correct.

- This is a **safety** property.

- **Deadlock-freedom property:**

- If multiple threads simultaneously request to enter a critical section, then it must allow one to proceed.

- Threads outside the critical section have no say in which thread can proceed into the critical section, only those currently waiting have influence.
- It implies the system never “freezes”.
- This is a **liveness** property.

Formalizing Critical Sections (cont.)

- Starvation-freedom property:

- Every thread that attempts to acquire the lock eventually succeeds.
- Many mutual exclusive algorithms in practice are not starvation free because its less likely starvation will occur in those algorithms.
- There is no guarantee on how long thread will wait to acquire the lock.
- Also known as **lockout freedom** or **bounded-waiting**
- This is a **liveness** property.

- Fairness Principle:

- A thread who just left the critical section cannot immediately re-enter the critical section if other threads have already requested to enter the critical section.
- Some algorithms place bounds on how long a thread can wait.

Assignment Project Exam Help

Theoretical ways of implementing locks:

- See Companion Slides: 2-44

Add WeChat powcoder

Assignment Project Exam Help

Parallel Performance
<https://powcoder.com>

Add WeChat powcoder

Amdahl's Law

To calculate how much a computation can be sped up by running part of it in parallel, can be done using **Amdahl's Law**:

Assignment Project Exam Help

- Potential program speedup is defined by the fraction of code (P) that can be parallelized

$$\text{Speedup} = \frac{1}{(1 - P)} \quad (1)$$

<https://powcoder.com>

- If none of the code can be parallelized, $P = 0$ and the speedup = 1 (no speedup).
- If all of the code is parallelized, $P = 1$ and the speedup is infinite (in theory). If 50% of the code can be parallelized, maximum speedup = 2, meaning the code will run twice as fast.
- See **Companion slides: 34-53** for more examples.

Amdahl's Law with Processors

With adding in the number of processors, the speed up can be modeled as such:

Assignment Project Exam Help

$$\text{Speedup} = \frac{1}{\frac{P}{N} + S}$$

P = parallel fraction, N = number of processors, and S = serial fraction

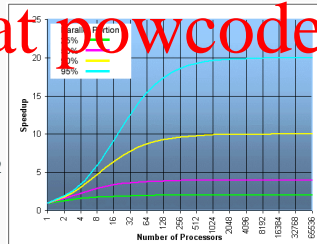
<https://powcoder.com>

It becomes obvious that there are limits to the scalability of parallelism:

Add WeChat powcoder

speedup				
N	P = .50	P = .90	P = .95	P = .99
10	1.82	5.26	6.89	9.17
100	1.98	9.17	16.80	50.25
1,000	1.99	9.91	19.62	90.99
10,000	1.99	9.91	19.96	99.02
100,000	1.99	9.99	19.99	99.90

2



3

Assignment Project Exam Help

Practical Implementation of Locks
<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Go has a package called `sync` that provides basic synchronization primitives such as mutual exclusion locks.

The `sync` package Go's provides mutual exclusion with `sync.Mutex` and its two methods:

- `m.Lock()`: locks m. If the lock is already in use, the calling goroutine blocks until the mutex is available.
- `m.Unlock()`: Unlocks m. It is a run-time error if m is not locked on entry to `Unlock`.

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

A lock requires making threads wait until it is their turn to enter the critical section. One implementation of a lock uses the notion of **spinning** to make a thread wait:

- When spinning, a thread repeatedly tests a condition, but, effectively, does no useful work until the condition has the appropriate value.
- Can be very wasteful of CPU cycles.
- Can also be unreliable if compiler optimization is turned on.
- Can be dangerous if not done in a deterministic way.

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

With caches, spinning becomes practical

- First time - Load flag bit into cache
- As long as it doesn't change
 - Hit in cache (no interconnect used)
- When it changes

One-time cost (a cache-miss)
Observes data has changed and must go grab it from main memory, which stops it from spinning.

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Many of the low-level synchronization primitives (e.g., locks, monitors, etc.) are built off of specialized hardware primitives/instructions (also known as **atomic operations**):

- On a shared memory system, an operation is considered **atomic** if it completes in a single step relative to other threads.
- No other thread can observe the modification to that shared variable half-way through its operation.

<https://powcoder.com>

Add WeChat powcoder

Review: Overview of Synchronization Primitives

Hardware provides simple low-level atomic operations:

- x86 load and store of words
- Special instructions:
 - compare-and-swap (CAS) (AMD, Intel, Sun)
 - pair of test-and-set instructions: loaded-linked and store-conditional (LL/SC) (ARM, IBM PowerPC, etc.)
- Provided in Go: `import "sync/atomic"`

We use those simple low-level atomic operations to build higher-level synchronization primitives:

- Lock (Today)
- Monitor
- Semaphore
- Conditional Variable
- Barrier

Problems with Atomic Operations

Best practices is to use atomic operations sparingly because:

Problem with atomic operations:

Most atomic operations are implemented using CAS or (LL/SC), which take significant more cycles to complete than a simple load or store instruction.

- Causes a memory fence which forces the write-back buffer to be sent to main memory. This process can then stall other processors from reading/writing to main memory
- Prevents out-of-order execution and various compiler optimizations.
- Cost to performance varies depending on architectures, program design, etc.
- Adds more hardware complexity.

Implementations

TTAS Lock

- TTAS Lock
- Exponential Backoff Lock
- Anderson Lock (ALock)
- CLH Lock
- MCS Lock
- Queue Locks
- Time-out Locks
- See Companion Slides: 54-242

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder